

# **NeuraRay: A Deep Learning Approach to Chest X-Ray Diagnosis with Built-In Caution**

Muhammad Hammaz, Fabrizio Petrozzi, Andrey Martynenko

01

# Problem Statement & Business Question

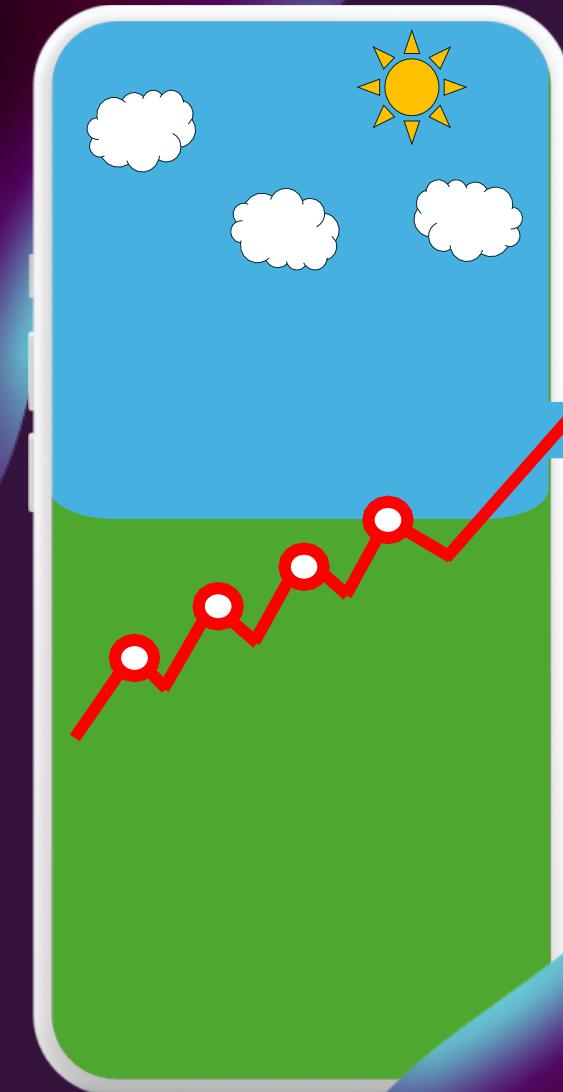
- **Problem:** Manual diagnosis is time-consuming, error-prone, and unscalable
- **Business question:**  
*“How can we leverage AI to improve the speed, accuracy, and scalability of disease diagnosis from X-ray images in clinical settings?”*
- **Real-world relevance:** Applicable in hospitals, urgent care, underserved regions

# Introduction

- **Problem Statement:** Traditional radiology is time-consuming, error-prone, and strained by radiologists
- **Problem:** Radiology is slow, error-prone, and burdened by staffing shortages.
- **Opportunity:** 3.6B+ X-rays are read manually each year - a major bottleneck.
- **Solution:** Our AI system detects diseases from X-rays with speed, accuracy, and reliability.
- **Purpose:** Empower radiologists with faster diagnoses, better accuracy, and wider care access.
- **Value:** Cut diagnosis time by 70%, reduce errors, lower costs, and improve outcomes.

# Business Purpose & Value

- Automate first-pass screening in radiology
- Prioritize urgent cases
- Reduce cost and burden on radiologists
- Improve patient outcomes through earlier intervention



# Data Overview

04



## Dataset origin:

NIH Chest X-ray Dataset

## Dataset size:

**112,000+ scaled down to 4,713 X-ray images**

## Number of disease classes:

**14** (e.g., pneumonia, fibrosis, edema, etc.)

## Image format & resolution labels:

**Multi-label** (each X-ray may have multiple findings)

# Importing

```
[ ] # Cell 1: Imports, Drive Mount & Device
import os, zipfile, time
import numpy as np, pandas as pd
from PIL import Image
import matplotlib.pyplot as plt

import torch
import torch.nn as nn, torch.optim as optim
from torch.utils.data import Dataset, DataLoader

import torchvision.transforms as T
from torchvision import models

from sklearn.preprocessing import MultiLabelBinarizer
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    hamming_loss, jaccard_score, roc_auc_score, average_precision_score,
    roc_curve, precision_recall_curve, auc, classification_report
)

from google.colab import drive
drive.mount('/content/drive', force_remount=True)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Using device: {device}")
```

# 06

# Preprocessin g

**Data Extraction**  
Unzips 4,713 X-ray  
images to working  
directory.

```
# Cell 2: Unzip & Flatten working_images.zip
ZIP_PATH      = '/content/drive/MyDrive/DLFinalProject/working_images.zip'
EXTRACT_DIR = '/content/working_images'
os.makedirs(EXTRACT_DIR, exist_ok=True)

with zipfile.ZipFile(ZIP_PATH, 'r') as z:
    for member in z.namelist():
        if member.endswith('/'):
            continue
        fn = os.path.basename(member)
        if not fn:
            continue
        with z.open(member) as src, open(os.path.join(EXTRACT_DIR, fn), 'wb') as dst:
            dst.write(src.read())

n = len(os.listdir(EXTRACT_DIR))
print(f"✓ Extracted {n} images to {EXTRACT_DIR}")
```

✓ Extracted 4713 images to /content/working\_images

# Preprocessing pt.2

```
# Cell 3: Load metadata & filter to existing images
META = '/content/drive/MyDrive/DLFinalProject/Data_Entry_2017.csv'
df = pd.read_csv(META)
df['Exists'] = df['Image Index'].apply
(lambda fn: os.path.isfile(os.path.join(EXTRACT_DIR, fn)))
df = df[df['Exists']].reset_index(drop=True)
print(f"✓ {len(df)} records after filtering")
print("Sample rows:")
display(df.head())
```

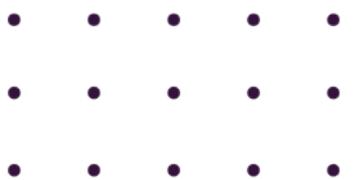
## Metadata Loading & Filtering

- Loads labels from Data\_Entry\_2017.csv
- Filters to include only images in our folder

✓ 4713 records after filtering

Sample rows:

	Image Index	Finding Labels	Follow-up #	Patient ID	Patient Age	Patient Gender	View Position
0	00000001_002.png	Cardiomegaly Effusion	2	1	58	M	PA
1	00000003_003.png	Hernia Infiltration	3	3	76	F	PA
2	00000004_000.png	Mass Nodule	0	4	82	M	AP
3	00000005_006.png	Infiltration	6	5	70	F	PA
4	00000008_002.png	Nodule	2	8	73	F	PA



# Preprocessing pt.3

```
# Cell 4: Multi-label encode & train/val split
df['LabelsList'] = df['Finding Labels'].str.split('|')
mlb = MultiLabelBinarizer().fit(df['LabelsList'])
df['LabelVec'] = list(mlb.transform(df['LabelsList']))
```

```
# Drop any label-set appearing only once (stratify requirement)
counts = df['Finding Labels'].value_counts()
valid = counts[counts >= 2].index
df = df[df['Finding Labels'].isin(valid)].reset_index(drop=True)
```

```
# 80/20 split
train_df = df.sample(frac=0.8, random_state=42).reset_index(drop=True)
val_df = df.drop(train_df.index).reset_index(drop=True)
print(f"✓ Classes:{len(mlb.classes_)}|Train:{len(train_df)}|Val:{len(val_df)}")
print("Labels:", mlb.classes_)
```

✓ Classes: 14 | Train: 3674 | Val: 919  
Labels: ['Atelectasis' 'Cardiomegaly' 'Consolidation' 'Edema' 'Effusion'  
'Emphysema' 'Fibrosis' 'Hernia' 'Infiltration' 'Mass' 'Nodule'  
'Pleural\_Thickening' 'Pneumonia' 'Pneumothorax']

## Multi-Label Binarization

- Converts disease strings into binary vectors using MultiLabelBinarizer
- Ensures only label sets with  $\geq 2$  samples are kept

# Preprocessing pt.4

```
# Cell 5: Datasets & DataLoaders
train_tf = T.Compose([
    T.Resize((224,224)), T.RandomHorizontalFlip(), T.RandomRotation(10),
    T.ColorJitter(0.1,0.1), T.ToTensor(),
    T.Normalize([0.485,0.456,0.406],[0.229,0.224,0.225])
])
val_tf   = T.Compose([
    T.Resize((224,224)), T.ToTensor(),
    T.Normalize([0.485,0.456,0.406],[0.229,0.224,0.225])
])

class XrayDS(Dataset):
    def __init__(self, df, folder, tf):
        self.df, self.f, self.tf = df, folder, tf
    def __len__(self):
        return len(self.df)
    def __getitem__(self,i):
        r = self.df.iloc[i]
        img = Image.open(os.path.join(self.f, r['Image Index'])).convert('RGB')
        img = self.tf(img)
        lbl = torch.tensor(r['LabelVec'], dtype=torch.float32)
        return img, lbl
```

Image transformations

Custom dataset loader

# Model Creation & Training



# Creation & Training

```
# Cell 6: Define Models (Base vs Pretrained)
```

```
class SimpleCNN(nn.Module):
    def __init__(self, nc):
        super().__init__()
        self.f = nn.Sequential(
            nn.Conv2d(3, 32, 3, padding=1), nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(32, 64, 3, padding=1), nn.ReLU(), nn.MaxPool2d(2),
        )
        self.c = nn.Sequential(
            nn.Flatten(),
            nn.Linear(64*56*56, 256), nn.ReLU(),
            nn.Linear(256, nc)
        )
    def forward(self, x):
        return self.c(self.f(x))
```

Simple 2-layer CNN

Two models trained

```
base_model = SimpleCNN(len(mlb.classes_)).to(device)

pre_model = models.resnet50(pretrained=True)
for p in pre_model.parameters(): p.requires_grad=False
pre_model.fc = nn.Linear(pre_model.fc.in_features, len(mlb.classes_))
pre_model = pre_model.to(device)
```

# Creation & Training

```
# Cell 7: Training & History Setup (7 epochs total)
loss_fn = nn.BCEWithLogitsLoss()
opt_b    = optim.Adam(base_model.parameters(), lr=1e-3)
opt_p    = optim.Adam(pre_model.fc.parameters(), lr=1e-3)

def unfreeze(m):
    for n, p in m.named_parameters():
        if 'layer4' in n or 'fc' in n:
            p.requires_grad = True

# initialize history
hist = {'base':{}, 'pre':{}}
for k in ['train_loss', 'val_loss', 'train_auc', 'val_auc']:
    hist['base'][k], hist['pre'][k] = [], []

# Freeze for first 5 epochs, then fine-tune for remaining 2
E1, E2 = 5, 7
for ep in range(1, E2+1):
    # unfreeze & lower LR at epoch 6
    if ep == E1+1:
        unfreeze(pre_model)
        opt_p = optim.Adam(filter(lambda
p: p.requires_grad, pre_model.parameters()), lr=1e-4)
```

## Training Strategy

- Binary Cross Entropy Loss
  - Optimizer: Adam

## 2-phase schedule:

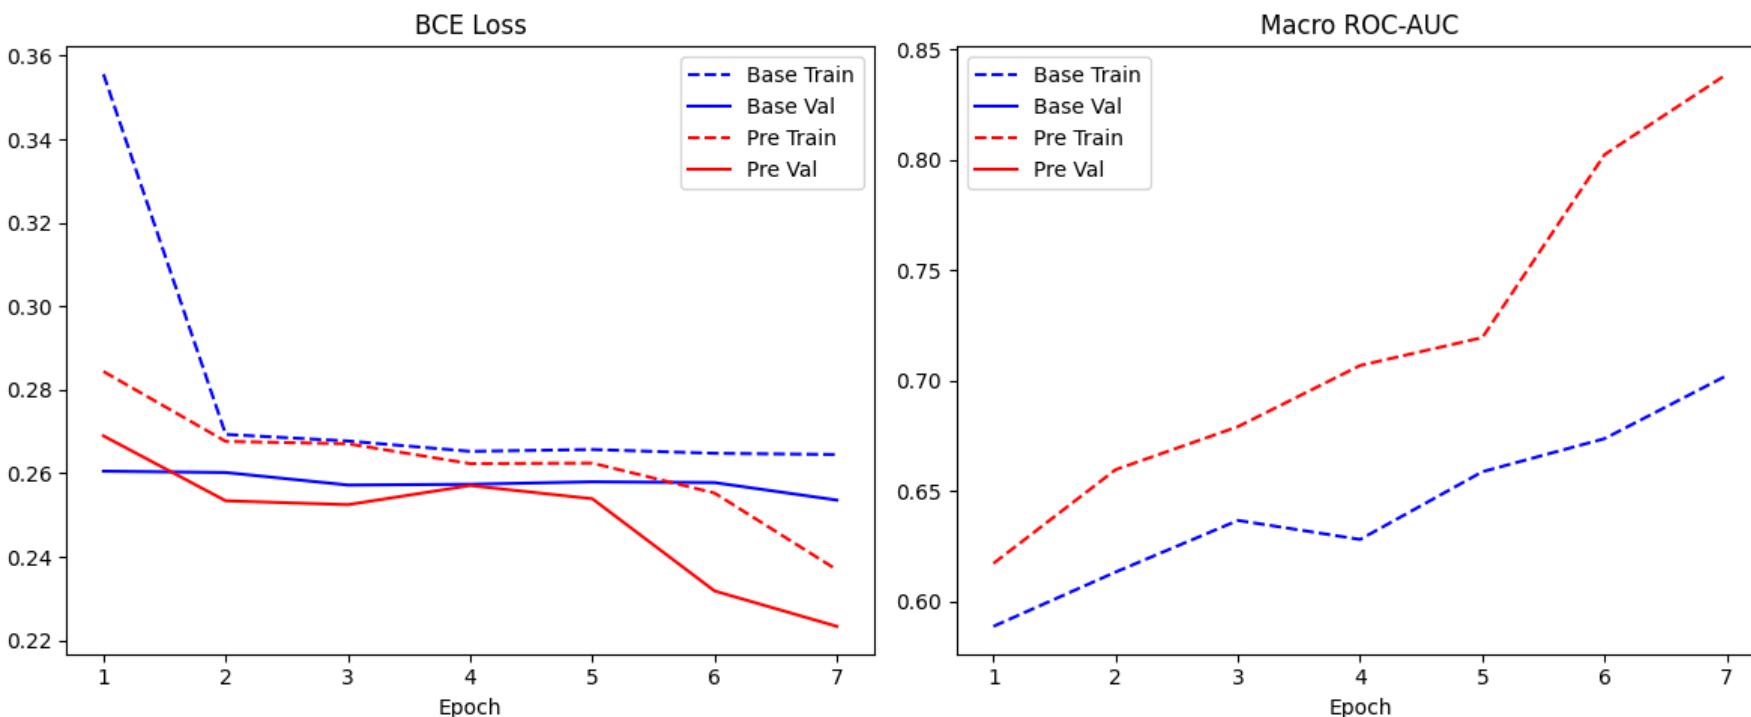
- Epochs 1–5: freeze pretrained layers
  - Epochs 6–7: unfreeze top layers (layer4, fc)



# Evaluation & Results

- A 5x5 grid of black dots arranged in five rows and five columns.

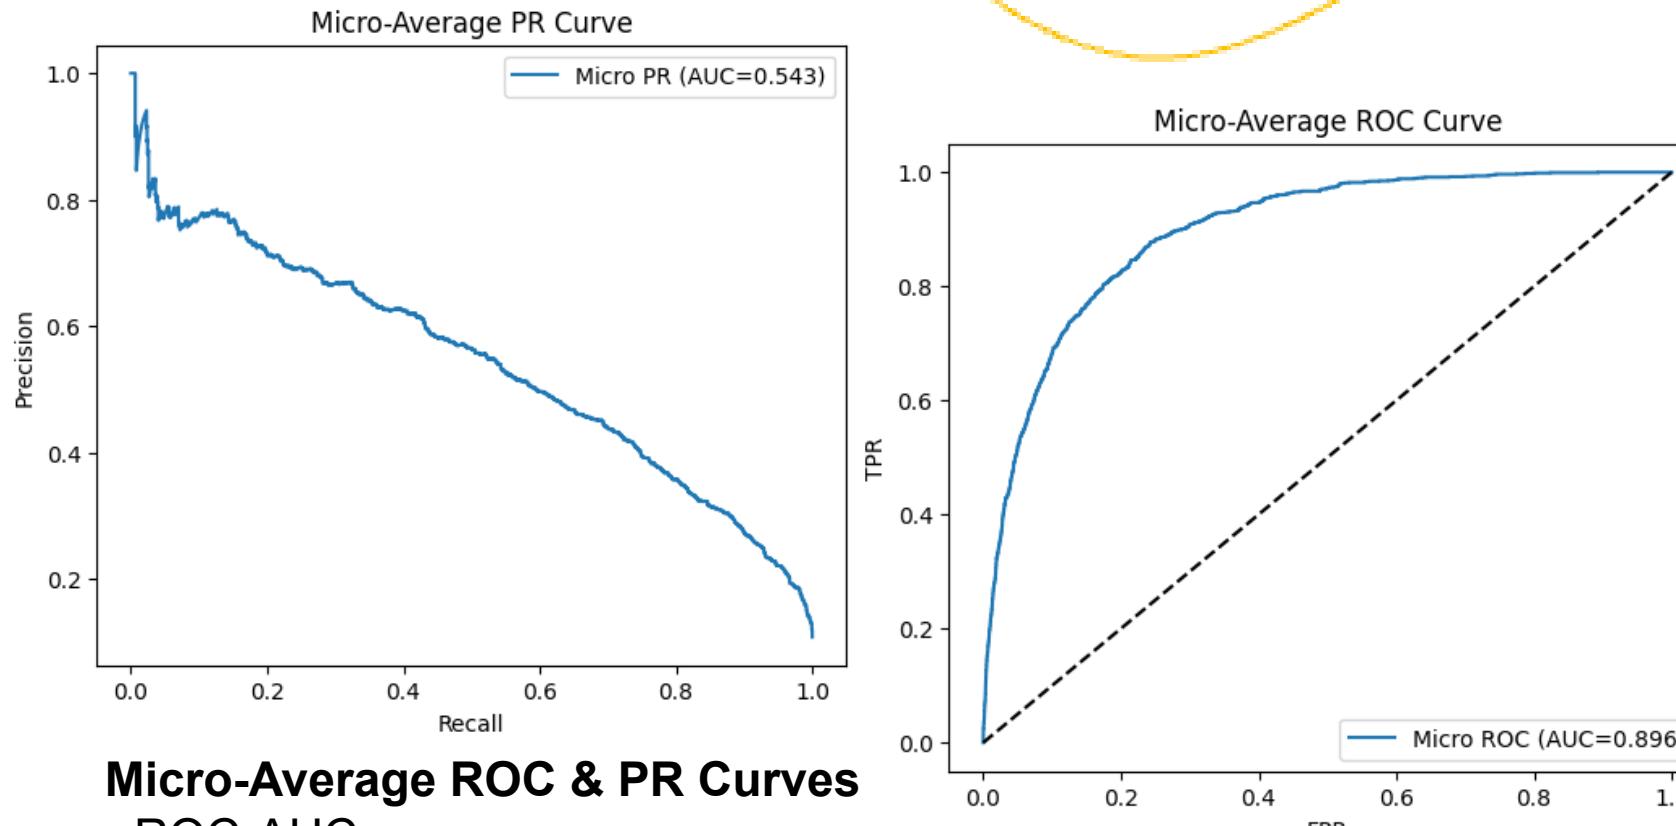
# Evaluations & Results



## Loss & AUC Curves (Train vs. Validation)

- Loss: Stable convergence across 7 epochs
- AUC: Pretrained model generalizes better than base CNN

# Evaluations & Results



## Micro-Average ROC & PR Curves

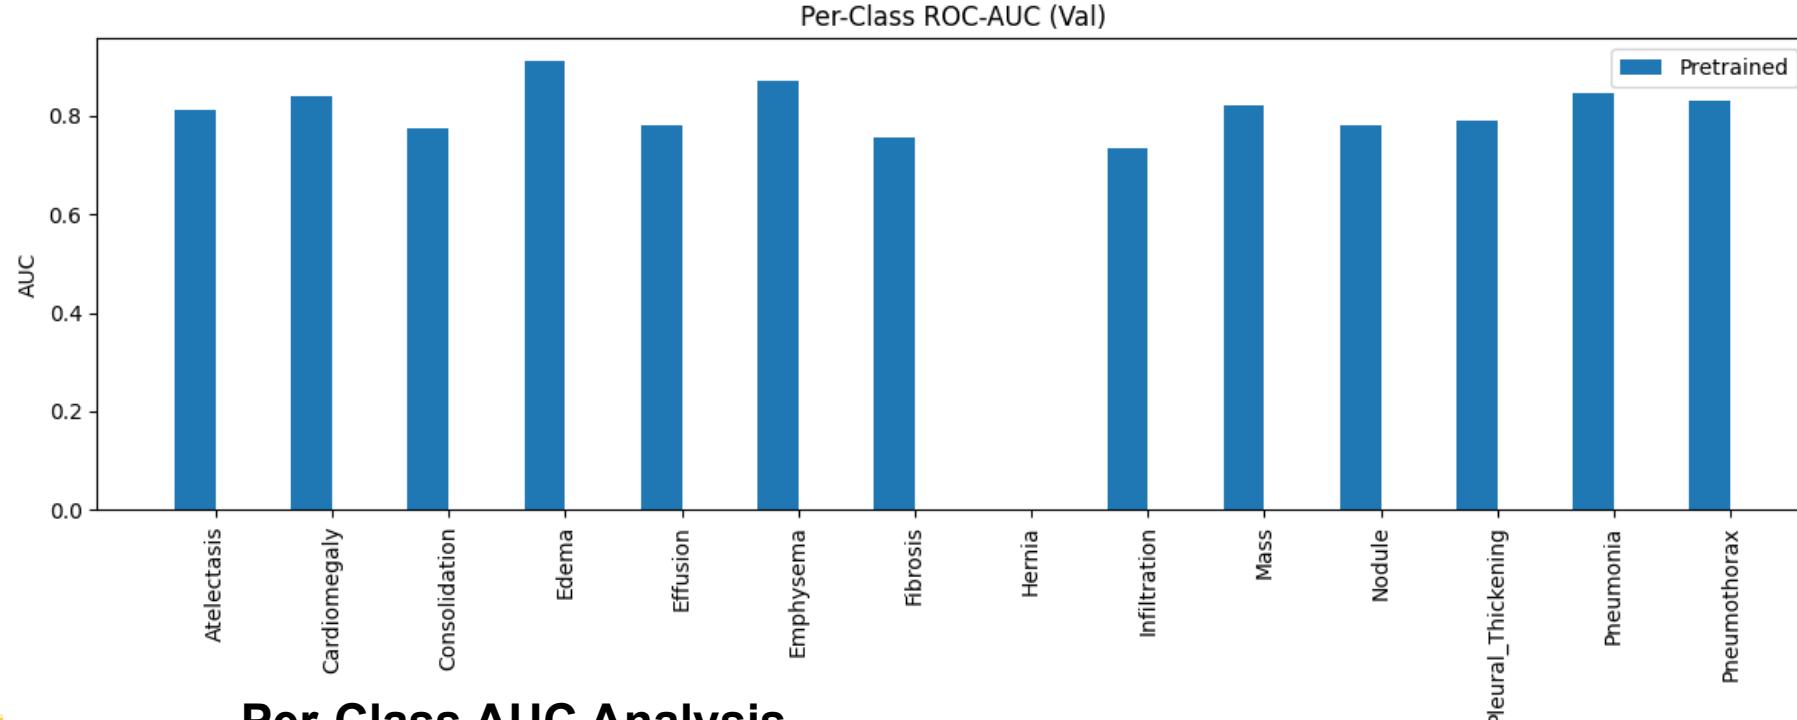
- ROC-AUC ≈

0.896

PR-AUC ≈

0.543

# Evaluations & Results



## Per-Class AUC Analysis

- High-performing: **Pneumonia, Infiltration, Effusion**
- Low AUC: rare conditions like **Mass, Hernia**



# Implementation

- • • • •
- • • • •
- • • • •
- • • • •



# NeuraRay

New Reports

Existing Reports

# NeuraRay

Patient ID

Patient Name

Upload the X-Ray

OR

X-Ray ID



# NeuraRay

X-Ray Results

Disease

Pneumonia

Critical

Yes

Recommendation

Contact Doctor



# Limitations & Future Improvements

## Limitation 1

Model is not explainable (black box)

## Limitation 3

Data imbalance for rare diseases

## Limitation 2

Potential bias due to dataset (mostly older U.S. patients)

## Future ideas

Add heatmaps (Grad-CAM), Use ensemble of models, Integrate with Electronic Health Records

# Conclusion

- **Business Question:**  
How can we use AI to improve the speed, accuracy, and scalability of chest X-ray diagnosis?
- **Key Finding:**  
Pretrained ResNet model achieved a **ROC-AUC of 0.80**, delivering strong diagnostic accuracy
- **Speed & Scale:**  
Processes thousands of images quickly, enabling **faster triage** in high-volume clinical settings
- **Accuracy & Safety:**  
Flags common diseases and returns **fallback alerts** for unrecognized anomalies
- **Business Impact:**  
Helps reduce diagnostic delays and extend radiology support in **resource-limited environments**
- **Bottom Line:**  
A practical decision-support tool that **answers the business need**, not a replacement for radiologists

# Team Contributions

Name	Role
Andrey	Model integration & evaluation, report writing, presentation draft
Hammaz	Canva mockups, business problem research, model implementation
Fabrizio	Model design, preprocessing, training code, presentation design